

IAML: Decision Trees

Victor Lavrenko and Charles Sutton
School of Informatics

Semester 1

Predict if John will play tennis

Training examples: 9 yes / 5 no

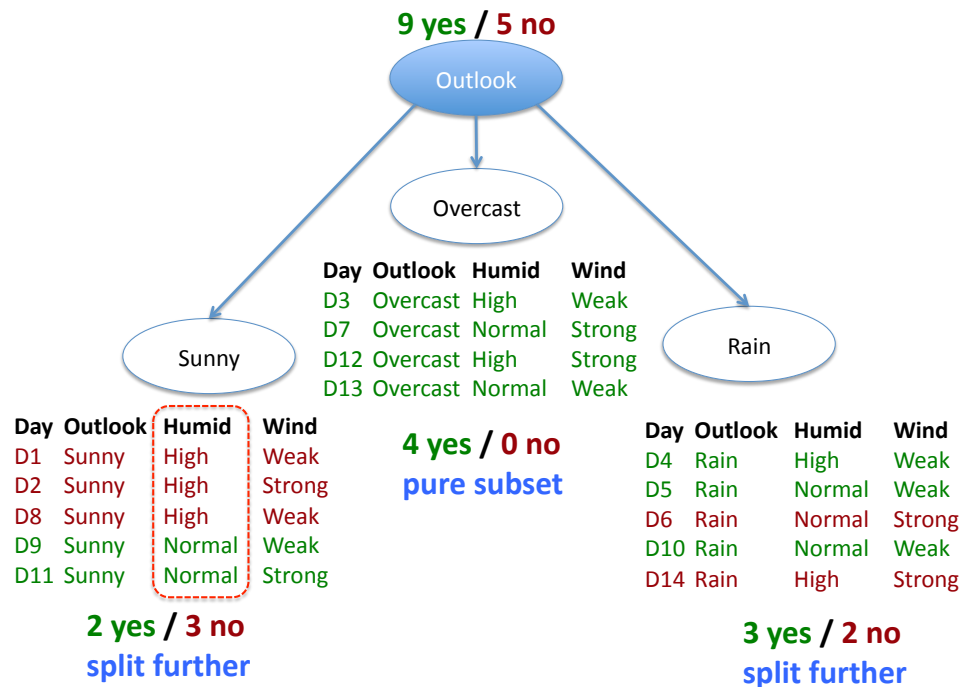
- Hard to guess
- Divide & conquer:
 - split into subsets
 - are they pure? (all yes or all no)
 - if yes: stop
 - if not: repeat
- See which subset new data falls into

Day	Outlook	Humidity	Wind	Play
D1	Sunny	High	Weak	No
D2	Sunny	High	Strong	No
D3	Overcast	High	Weak	Yes
D4	Rain	High	Weak	Yes
D5	Rain	Normal	Weak	Yes
D6	Rain	Normal	Strong	No
D7	Overcast	Normal	Strong	Yes
D8	Sunny	High	Weak	No
D9	Sunny	Normal	Weak	Yes
D10	Rain	Normal	Weak	Yes
D11	Sunny	Normal	Strong	Yes
D12	Overcast	High	Strong	Yes
D13	Overcast	Normal	Weak	Yes
D14	Rain	High	Strong	No

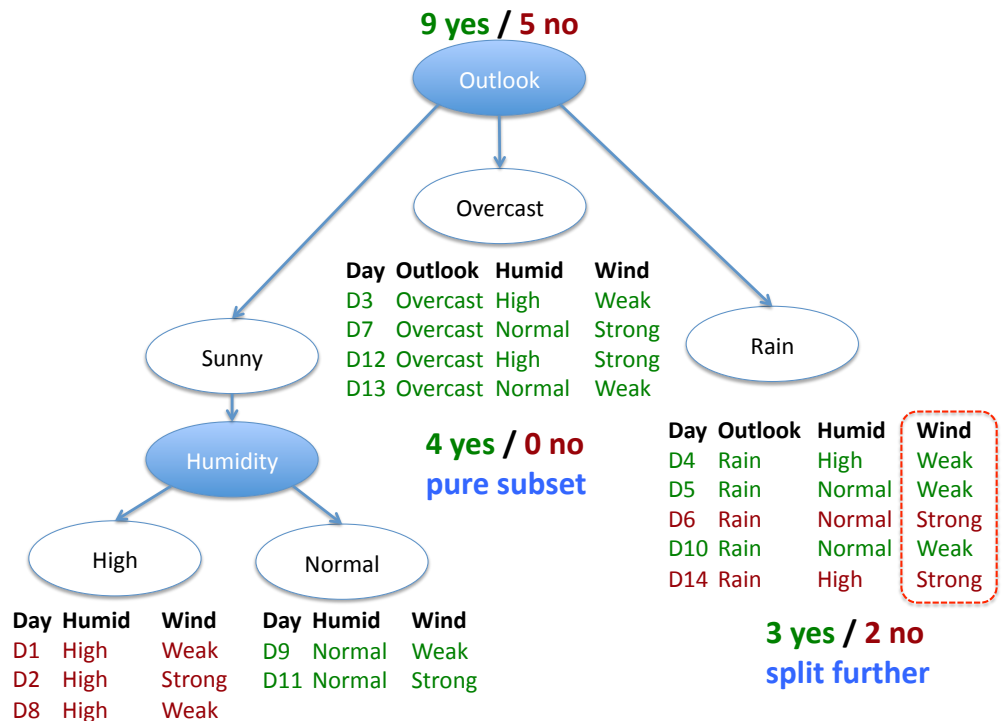
New data:

D15	Rain	High	Weak	?
-----	------	------	------	---

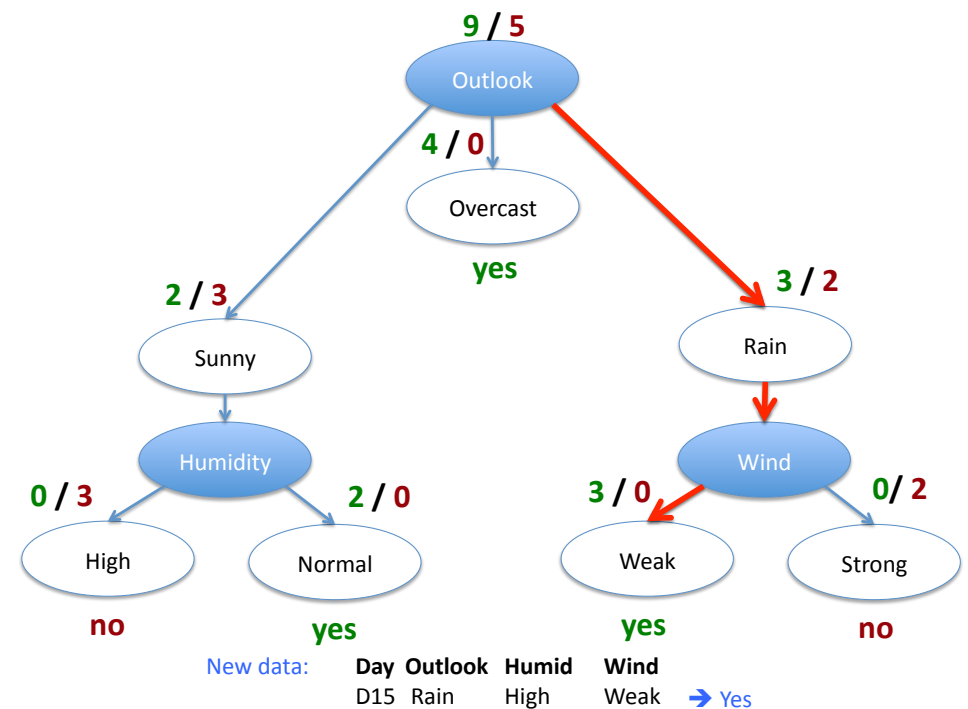
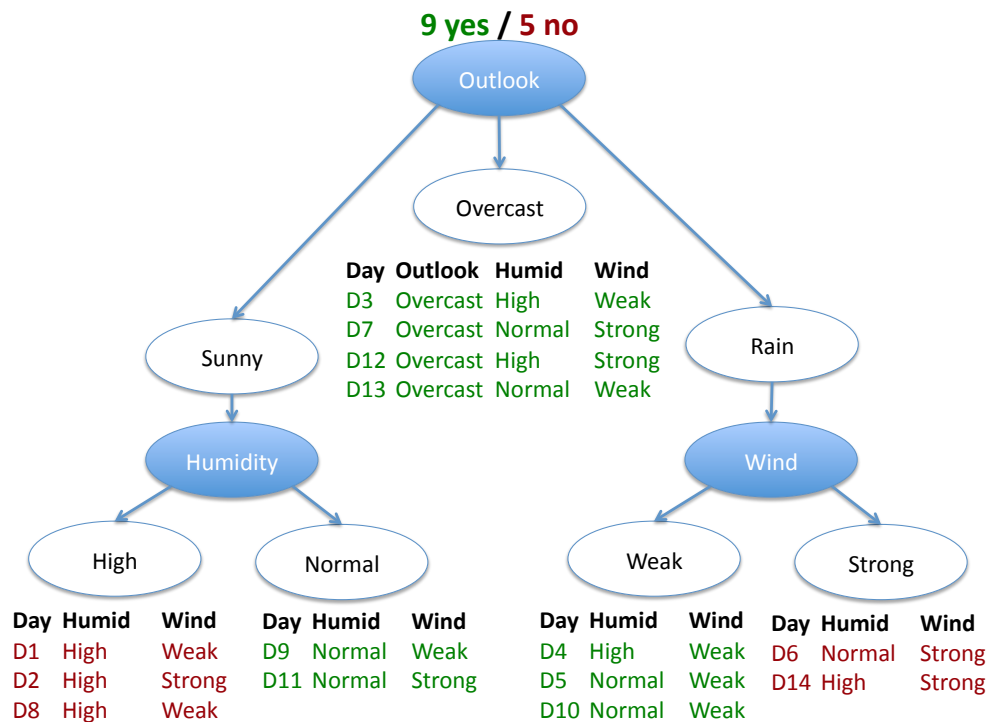
Copyright © 2011 Victor Lavrenko



Copyright © 2011 Victor Lavrenko




Copyright © 2011 Victor Lavrenko



ID3 algorithm

- Split (node, {examples}):
 - $A \leftarrow$ the **best attribute** for splitting the {examples}
 - Decision attribute for this node $\leftarrow A$
 - For each value of A, create new child node
 - Split training {examples} to child nodes
 - If examples perfectly classified: STOP
else: iterate over new child nodes
Split (child_node, {subset of examples})
- Ross Quinlan (ID3: 1986), (C4.5: 1993)
- Breimanetal (CaRT: 1984) from statistics

Which attribute to split on?



- Want to measure “purity” of the split
 - more certain about Yes/No after the split
 - pure set (4 yes / 0 no) => completely certain (100%)
 - impure (3 yes / 3 no) => completely uncertain (50%)
 - can’t use $P(\text{“yes”} \mid \text{set})$:
 - must be symmetric: 4 yes / 0 no as pure as 0 yes / 4 no

Entropy


- Entropy: $H(S) = -p_{(+)} \log_2 p_{(+)} - p_{(-)} \log_2 p_{(-)}$ bits
 - S ... subset of training examples
 - $p_{(+)} / p_{(-)}$... % of positive / negative examples in S
- Interpretation: assume item X belongs to S
 - how many bits need to tell if X positive or negative

- impure (3 yes / 3 no):

$$H(S) = -\frac{3}{6} \log_2 \frac{3}{6} - \frac{3}{6} \log_2 \frac{3}{6} = 1 \text{ bits}$$

- pure set (4 yes / 0 no):

$$H(S) = -\frac{4}{4} \log_2 \frac{4}{4} - \frac{0}{4} \log_2 \frac{0}{4} = 0 \text{ bits}$$



Copyright © 2011 Victor Lavrenko

Information Gain

- Want many items in pure sets
- Expected drop in entropy after split:


$$Gain(S, A) = H(S) - \sum_{V \in \text{Values}(A)} \frac{|S_V|}{|S|} H(S_V)$$

V ... possible values of A
 S ... set of examples $\{X\}$
 S_V ... subset where $X_A = V$

- Mutual Information

– between attribute A and class labels of S


$$\begin{aligned} \text{Gain}(S, \text{Wind}) &= H(S) - \frac{8}{14} H(S_{\text{weak}}) - \frac{6}{14} H(S_{\text{strong}}) \\ &= 0.94 - \frac{8}{14} * 0.81 - \frac{6}{14} * 1.0 \\ &= 0.049 \end{aligned}$$



Copyright © 2011 Victor Lavrenko

Overfitting in Decision Trees

- Can always classify training examples perfectly
 - keep splitting until each node contains 1 example
 - singleton = pure
- Doesn't work on new data



Copyright © 2011 Victor Lavrenko

Figure credit: Tom Mitchell, 1997

Avoid overfitting

- Stop splitting when not statistically significant
- Grow, then post-prune
 - based on validation set
- Sub-tree replacement pruning (WF 6.1)
 - for each node:
 - pretend remove node + all children from the tree
 - measure performance on validation set
 - remove node that results in greatest improvement
 - repeat until further pruning is harmful

Copyright © 2011 Victor Lavrenko


General Structure

- **Task:** classification, discriminative
- **Model structure:** decision tree
- **Score function**
 - information gain at each node
 - preference for short trees
 - preference for high-gain attributes near the root
- **Optimization / search method**
 - greedy search from simple to complex
 - guided by information gain

Copyright © 2011 Victor Lavrenko

Problems with Information Gain

- Biased towards attributes with many values



- Won't work for new data: D15 Rain High Weak

- Use GainRatio:

$$SplitEntropy(S,A) = - \sum_{V \in Values(A)} \frac{|S_V|}{|S|} \log \frac{|S_V|}{|S|}$$

A ... candidate attribute
V ... possible values of A
S ... set of examples {X}
S_V ... subset where X_A = V

$$GainRatio(S,A) = \frac{Gain(S,A)}{SplitEntropy(S,A)}$$


penalizes attributes with many values

Copyright © 2011 Victor Lavrenko

Trees are interpretable

- Read rules off the tree
 - concise description of what makes an item positive

- No “black box”
 - important for users



Rule:
(Outlook = Overcast) V
(Outlook = Rain ∧ Wind = Weak) V
(Outlook = Sunny ∧ Humidity = Normal)

Copyright © 2011 Victor Lavrenko

Figure credit: Tom Mitchell, 1997

Continuous Attributes

- Dealing with continuous-valued attributes:
 - create a split: (Temperature > 72.3) = True,False
- Threshold can be optimized (WF 6.1)




Figure credit: Chris Bishop, PRML

Multi-class and Regression

- Multi-class classification:
 - predict most frequent class in the subset
 - entropy: $H(S) = - \sum_c p_{(c)} \log_2 p_{(c)}$
 - $p_{(c)}$... % of examples of class c in S
- Regression:
 - predicted output = mean of the training examples in the subset
 - requires a different definition of entropy
 - can use linear regression at the leaves (WF 6.5)

Copyright © 2011 Victor Lavrenko

Summary

- ID3: grows decision tree from the root down
 - greedily selects next best attribute (Gain)
- Searches a complete hypothesis space
 - prefers smaller trees, high gain at the root
- Overfitting addressed by post-pruning
 - prune nodes, while accuracy \uparrow on validation set
- Can handle missing data (see WF 6.1)
- Easily handles irrelevant variables
 - Information Gain = 0 => will not be selected

Copyright © 2011 Victor Lavrenko

Random Decision Forest

- Grow K different decision trees:
 - pick a random subset S_r of training examples
 - grow a full ID3 tree (no pruning):
 - when splitting: pick from $d \ll D$ random attributes
 - compute gain based on S_r instead of full set
 - repeat for $r = 1 \dots K$
- Given a new data point X :
 - classify X using each of the K trees
 - use majority vote: class predicted most often
- Fast, scalable, state-of-the-art performance

Copyright © 2011 Victor Lavrenko